Conditioning and Non-Convexity of Semantic Decoding by Inversion of Deep Brain Encoding Models

Submitted in partial fulfillment of the requirements for the degree of

Bachelor of Technology
in
Electronics and Communication Engineering (specialisation in Biomedical Engineering)

by

ABHRANIL SARKAR (23BML0077)

TANISHA CHOUDHURI (23BML0081)

Under the guidance of Dr. Sharmila Nageswaran

School of Electronics Engineering (SENSE), VIT, Vellore.



October, 2026

DECLARATION

We hereby declare that the thesis entitled “Conditioning and Non-Convexity of Semantic Decoding by Inversion of Deep Brain Encoding Models” submitted by us, for the award of the degree of Bachelor of Technology in Electronics and Communication Engineering (specialisation in Biomedical Engineering) to VIT is a record of bonafide work carried out by us under the supervision of Dr. Sharmila Nageswaran.

We further declare that the work reported in this thesis has not been submitted previously to this institute or anywhere for the consideration of the degree/diploma.

Place : Vellore

Date :

Signature of the Candidates

CERTIFICATE

This is to certify that the thesis entitled “Conditioning and Non-Convexity of Semantic Decoding by Inversion of Deep Brain Encoding Models” submitted by Abhranil Sarkar (23BML0077) and Tanisha Choudhuri (23BML0081), School of Electronics Engineering (SENSE), VIT, for the award of the degree of Bachelor of Technology in Electronics and Communication Engineering (specialisation in Biomedical Engineering), is a record of bonafide work carried out by them under my supervision during the period July 2026 to October 2026, as per the VIT code of academic and research ethics.

The contents of this report have not been submitted either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university. The thesis fulfills the requirements and regulations of the University and in my opinion meets the necessary standards for submission.

Place : Vellore

Date : **Signature of the Guide**

Internal Examiner

ACKNOWLEDGEMENTS

We express our sincere gratitude to our guide, Dr. Sharmila Nageswaran, Associate Professor, School of Electronics Engineering (SENSE), VIT Vellore, for her guidance, encouragement and critical feedback at every stage of this project.

We thank the Dean and the Head of the Department, School of Electronics Engineering, and the members of the review panels for their suggestions, which shaped the final form of the proposed work, in particular the questions on the data behind the TRIBE v2 model and on the choice of fMRI over EEG.

We thank Meta FAIR for releasing the TRIBE v2 model and code openly, the authors of the LeBel et al. natural-language fMRI dataset for making their data public on OpenNeuro, and the developers of the open-source tools on which the proposed work depends.

Finally, we thank our families and friends for their constant support.

Abhranil Sarkar **Tanisha Choudhuri**

Executive Summary

Clinical problem

People with amyotrophic lateral sclerosis (ALS), brainstem stroke, locked-in syndrome or severe aphasia often keep their understanding of language but lose speech and writing. A non-invasive decoder that reads meaning from brain activity would give them a communication channel without neurosurgery. Such decoders are built on brain encoding models, which now predict fMRI responses to natural stimuli with high accuracy, and reliable decoding is widely assumed to follow from accurate encoding. That assumption had not been tested.

Rationale for the approach

TRIBE v2, Meta's open-weight 177-million-parameter trimodal brain encoder, is the most accurate public fMRI encoder, so inverting it tests directly whether an accurate encoder is also invertible. fMRI was chosen over EEG for its millimetre-scale detail at 20,484 cortical points, because a foundation encoder exists only for fMRI, and because EEG would add a second ill-posed inverse problem (source localisation). A synthetic ground truth placed in the encoder's own input space makes the recovery error exactly measurable, which real fMRI alone cannot provide.

Approach to the problem

A differentiable entry point after feature aggregation was identified, and the 1,152-dimensional stimulus representation was recovered by gradient descent through the frozen network. The inverse was characterised with exact synthetic ground truth across optimisation budget, starting points, measurement noise, cortical extent, priors and six cortical networks, and was then validated on real fMRI from a story that TRIBE v2 never saw during training (LeBel et al., 2023).

Key results

Encoding error fell 28.4-fold while recovery saturated at cosine 0.944. Equally good fits from different starting points disagreed (minimum pairwise cosine 0.817), and forward-only tests showed barriers between solutions (median 4.8×) rather than a null space: the inverse is non-convex, and its errors lie in the operator's least-sensitive directions. Across six cortical networks, the most responsive (visual) was the hardest to decode from and the least responsive (somatomotor) the easiest. On real fMRI, inversion identified the heard segment in 4 of 5 cases among 45 (chance 2.2%) while recovering little of the representation (cosine 0.02).

Conclusion

Encoding accuracy does not certify decoding. Recoverability has to be measured directly, region by region, before an encoder is relied on for assistive communication.

CONTENTS

1. INTRODUCTION.....1
1.1 Literature Review.....1
1.2 Research Gap.....2
1.3 Problem Statement.....3
2. RESEARCH OBJECTIVE.....4
3. RELEVANCE OF PROBLEM STATEMENT W.R.T SDG.....5
4. PROPOSED SYSTEM.....6
4.1 Design Approach / Materials & Methods.....6
4.1.1 The forward operator: TRIBE v2.....6
4.1.2 Training data of TRIBE v2.....6
4.1.3 Choice of fMRI over EEG.....7
4.1.4 Inversion method.....8
4.2 Codes and Standards.....9
4.3 Constraints, Alternatives and Tradeoffs.....9
5. PROJECT DESCRIPTION.....10
5.1 Instrumentation of the encoder.....10
5.2 Synthetic evaluation with exact ground truth.....10
5.2.1 Justification and validity of the ground truth.....10
5.3 Discriminating the causes of incomplete recovery.....12
5.4 Regional analysis.....13
5.5 Real-fMRI validation.....13
6. HARDWARE/SOFTWARE TOOLS USED.....15
7. SCHEDULE AND MILESTONES.....16
8. RESULT ANALYSIS.....17
8.1 Encoding accuracy does not certify decoding.....17
8.2 Equally good fits give different answers.....17
8.3 Barriers, not a null space.....18
8.4 Errors lie in the operator's weakest directions.....20
8.5 Further conditioning results.....20
8.6 Predictable is not recoverable, region by region.....23
8.7 Real fMRI.....25
8.8 Discussion.....26
9. CONCLUSION.....27
9.1 Obtained Results.....27
9.2 Future improvement/work.....27
9.3 Individual contribution from team members.....28
10. SOCIAL AND ENVIRONMENTAL IMPACT.....29

11. COST ANALYSIS..... 30
12. PROJECT OUTCOME PUBLICATION/PATENT..... 31
13. REFERENCES..... 32
APPENDIX A: CODE AND REPRODUCTION..... 35

LIST OF FIGURES

Figure 4.1: Inversion pipeline: the network is re-entered after feature aggregation; only the representation x is optimised, the network stays frozen.....6

Figure 5.1: Synthetic ground truth and a real stimulus representation along their top eight principal directions, and their effective dimensionality, variance in eight directions and lag-1 autocorrelation (real: 246 windows of the six LeBel stories).....11

Figure 5.2: Inversion with a real stimulus representation as ground truth: recovery per start (left) and prediction error along the straight lines between solutions (right).....12

Figure 8.1: Encoding error against recovery cosine across independent step budgets.....17

Figure 8.2: Agreement among fit-matched solutions across four ground truths (left) and fit versus recovery per initialisation (right).....18

Figure 8.3: Left: prediction error along the straight line between fit-matched blind solutions (blue, ten pairs) and from the best blind solution to the near-truth solution (orange). Right: final error and recovery for each of the eight starts.....19

Figure 8.4: Directional sensitivity of the operator in the original run: solution-difference directions against 200 random directions (the corrected pairs give 2.15 ± 0.17× over four truths).....20

Figure 8.5: Output change under rescaling of the ground truth relative to the optimiser's noise floor.....21

Figure 8.6: Recovery, relative error and retrieval against measurement noise.....21

Figure 8.7: Recovery against the number of observed cortical vertices.....21

Figure 8.8: Effect of ℓ_2 and temporal-smoothness priors in the under-determined regime (500 vertices).....22

Figure 8.9: Gradient inversion against the best ridge decoder.....22

Figure 8.10: The six anatomical networks on the inflated fsaverage5 surface.....23

Figure 8.11: Forward-sensitivity share of the text-, audio- and video-derived dimensions across cortex.....23

Figure 8.12: Recovery from each network versus size-matched random vertex sets.....24

Figure 8.13: Forward sensitivity against recoverability across networks.....24

Figure 8.14: Participation ratio of each network's responses against size-matched controls.....24

Figure 8.15: Real fMRI, unseen test story: alignment (A), encoding accuracy by network (B) and identification accuracy (C).....25

LIST OF TABLES

Table 1.1: Summary of the literature most relevant to this project.....2
Table 1.2: Novelty: closest prior work against the four properties combined here (● yes, ○ no).....3
Table 3.1: Mapping to specific SDG targets.....5
Table 4.1: TRIBE v2 training data: participants and stimuli.....7
Table 4.2: TRIBE v2 evaluation data (never used for training).....7
Table 4.3: fMRI versus EEG for this project.....8
Table 4.4: Standards and frameworks the project follows.....9
Table 4.5: Constraints and the trade-offs adopted.....9
Table 5.1: Evidence that the synthetic ground truth is valid.....12
Table 6.1: Hardware and software tools, their use and their cost to the project.....15
Table 7.1: Schedule and milestones.....16
Table 8.1: Encoding error and recovery across optimisation budgets.....17
Table 8.2: Agreement among fit-matched solutions, four ground truths, 14,000 steps.…18
Table 8.3: Straight-path barrier between fit-matched blind solutions: all ten pairs for ground truth 0, and the summary over four ground truths (14,000 steps).....19
Table 8.4: Further characterisation of the inverse (whole cortex unless stated).....22
Table 8.5: Network sensitivity and recoverability (synthetic, 8,000 steps).....25
Table 8.6: Real-fMRI validation on the unseen test story wheretheressmoke (UTS03).…25
Table 9.1: Individual contributions.....28
Table 11.1: Estimated and incurred cost.....30
Table A.1: Map from report results to code.....35

LIST OF ABBREVIATIONS

BCI	Brain–Computer Interface
BIDS	Brain Imaging Data Structure
BOLD	Blood-Oxygen-Level-Dependent (signal)
CER	Character Error Rate
CNeuroMod	Courtois Project on Neuronal Modelling
DMN	Default Mode Network
EEG	Electroencephalography
FAIR	Findable, Accessible, Interoperable, Reusable
fMRI	functional Magnetic Resonance Imaging
GPU	Graphics Processing Unit
HCP	Human Connectome Project
HRF	Haemodynamic Response Function
IEM	Inverted Encoding Model
LLM	Large Language Model
MEG	Magnetoencephalography
MNI	Montreal Neurological Institute (template space)
MSE	Mean Squared Error
PR	Participation Ratio
ROI	Region of Interest
SDG	Sustainable Development Goal
TR	Repetition Time
TRIBE	TRImodal Brain Encoder
WK / WP	Knowledge profile / Problem attribute (Washington Accord)

SYMBOLS AND NOTATIONS

x	Aggregated stimulus representation, x ∈ ℝ^(T×H), T = 32 (synthetic) or 200 (real), H = 1,152
$\hat{x}$	Representation recovered by inversion
$f(\cdot)$	Frozen TRIBE v2 forward operator from x to predicted fMRI
y	Observed or target cortical response, y ∈ ℝ^(V×T′), V = 20,484, T′ = 100
d	Difference between two fit-matched solutions
$s(v)$	Directional sensitivity of f along unit direction v
ε	Finite-difference step size
λ	Regularisation weight
r	Pearson correlation coefficient
α	Gain (rescaling) factor applied to x

1. INTRODUCTION

A brain encoding model predicts the brain response that a stimulus evokes; decoding runs the other way, from measured brain activity back to the stimulus. For fifteen years the dominant route to decoding has been through encoding: fit a forward model, then identify or reconstruct the stimulus whose predicted response best matches what was measured [1], [2]. Encoding models have since grown from linear maps on hand-built features [3], [4] to deep networks trained on hundreds of hours of fMRI, and their accuracy keeps improving with model size and data [5], [6].

Decoding meaning from non-invasive brain recordings matters clinically. Amyotrophic lateral sclerosis, brainstem stroke, locked-in syndrome and severe aphasia can leave comprehension intact while removing speech and writing. A non-invasive decoder would restore a communication channel without neurosurgery. It also raises ethical questions of mental privacy, so the limits of what can be decoded matter as much as the successes.

This project addresses a question that the field has largely skipped: whether an encoder that becomes very accurate also becomes invertible. A forward model can be highly predictive and still have an inverse that is non-unique or unstable [7], [8]. The question became answerable in 2026, when Meta released TRIBE v2, an open-weight foundation encoder that predicts fMRI at 20,484 cortical locations from video, audio and language [6].

1.1 Literature Review

Encoding models. Mitchell et al. predicted fMRI responses to nouns from word co-occurrence statistics [3]; Huth et al. mapped semantic selectivity across the whole cortex during natural speech [4]; and Antonello et al. showed that encoding accuracy with language-model features rises log-linearly with model size [5]. TRIBE won the Algonauts 2025 challenge [9], [10] and TRIBE v2 extends it to a foundation model trained on 451.6 h of fMRI and evaluated on 1,117.7 h from 720 subjects [6].

Decoding. Kay et al. identified which image a person was viewing by comparing measured responses with responses predicted for candidate images [2]. Tang et al. reconstructed continuous language from fMRI by searching word sequences with a language model and scoring them with an encoder [11]. With MEG and EEG, contrastive decoders identify perceived speech segments [12], and Brain2Qwerty decoded typed sentences with a 32% character error rate from MEG but 67% from EEG with the same architecture [13]. Shared-

subject models reduce the data each subject needs for image reconstruction [14]. All of these evaluate decoding by proxies such as rank, retrieval or word error, because the true representation behind real fMRI is unknown.

Inverting encoders. Inverted encoding models reconstruct channel responses from voxel patterns [15]. Gardner and Liu proved that such linear models recover the stimulus only up to an invertible linear transform [16]. In computer vision, inverting deep networks from their activations shows which information each layer keeps [17]. Two 2026 follow-ups use TRIBE v2: Bladon and Bent optimise inputs to maximise a region's predicted response (feature visualisation) [18], and Benchetrit et al. use TRIBE v2 forwards to synthesise training data for image decoders [19].

Table 1.1: Summary of the literature most relevant to this project

Work	Contribution	Limitation relevant here
Mitchell et al., 2008	Linear model predicts fMRI for unseen nouns	Single words, hand-built features
Kay et al., 2008	Identifies viewed images via predicted responses	Identification, not recovery
Naselaris et al., 2011	Encoding models as the route to decoding	Assumes invertibility
Huth et al., 2016	Semantic maps across cortex from natural speech	Linear; no inverse analysis
Antonello et al., 2023	Encoding accuracy scales with model size	Scaling shown for encoding only
Tang et al., 2023	Continuous language from fMRI via LM search	Per-subject decoder; search, not inversion
Gardner & Liu, 2019	Linear IEMs identifiable only up to a transform	Linear models only
Lévy et al., 2025	Typed sentences from MEG (CER 32%) and EEG (67%)	Shows EEG's limits for language
d'Ascoli et al., 2026 (TRIBE v2)	Open foundation encoder, 720 subjects	Encoder only; inverse never studied
Bladon & Bent; Benchetrit et al., 2026	Feature visualisation; synthetic training data	Neither tests uniqueness of the inverse

1.2 Research Gap

Encoders keep improving, decoders are trained separately and judged by proxies, and the inverse of a deep encoder has never been characterised: no study measures whether its solution is unique, how it is conditioned, or which parts of cortex carry recoverable information. A literature search on 9 October 2026 (arXiv, Google Scholar, general web; queries combining “TRIBE v2”, invert / inversion / identifiability / conditioning / non-convex and “brain encoding model”) found no published inverse or identifiability analysis of a deep foundation brain encoder. Table 1.2 places the closest work against the four properties this project

combines.

Table 1.2: Novelty: closest prior work against the four properties combined here (● yes, ○ no)

Study	Nonlinear deep encoder	Optimises encoder input	Tests uniqueness / conditioning	Region-wise recoverability
Kay et al., 2008	○	○	○	○
Gardner & Liu, 2019	○	●	●	○
Tang et al., 2023	○	●	○	●
TRIBE v2, 2026	●	○	○	○
Bladon & Bent, 2026	●	●	○	○
Benchetrit et al., 2026	●	○	○	○
Proposed work	●	●	●	●

1.3 Problem Statement

Characterise semantic decoding from fMRI as the inverse of a frozen, differentiable deep brain encoder (TRIBE v2): measure whether the inverse is unique, how it is conditioned, and where in cortex the stimulus representation is recoverable, first with exact synthetic ground truth and then on real naturalistic fMRI.

2. RESEARCH OBJECTIVE

1. Formulate semantic decoding as inversion of a frozen brain encoding operator and locate the point in the TRIBE v2 pipeline where the operator becomes differentiable.
2. Build a synthetic framework with exact ground truth, so that recovery error is measured exactly rather than through proxies.
3. Characterise the inverse with respect to optimisation budget, measurement noise, cortical extent, temporal rate of change and regularising priors.
4. Discriminate the causes of incomplete recovery: information loss (a null space), non-convexity, an optimiser artefact, or unidentifiable gain.
5. Benchmark gradient inversion against a linear decoder and map recoverability across six anatomical cortical networks.
6. Validate on real naturalistic fMRI (LeBel et al., 2023) with text and audio features, using a story that TRIBE v2 never saw during training.

All six objectives were achieved. Their evidence is presented in Chapter 8 and summarised in Chapter 9.

3. RELEVANCE OF PROBLEM STATEMENT W.R.T SDG

Table 3.1: Mapping to specific SDG targets

SDG target	Link to this project	Evidence from this project
SDG 3, Target 3.4 (promote mental health and well-being)	People with ALS, brainstem stroke or locked-in syndrome often keep comprehension but lose speech; non-invasive decoding could restore communication	Quantitative limits on the information fMRI decoding can recover set realistic clinical expectations
SDG 9, Target 9.5 (enhance research, upgrade technological capability)	Methodological contribution to neural-interface engineering with open models and data	Fully reproducible pipeline on free compute; public code, results and notebooks
SDG 10, Target 10.2 (inclusion irrespective of disability)	Communication access is a precondition for social and economic inclusion	Identifies which approaches are worth pursuing for assistive communication

The scope is stated honestly: this is enabling research, not a deployed assistive device. Its contribution to these targets is evidence about what non-invasive semantic decoding can achieve and what limits it.

4. PROPOSED SYSTEM

The proposed system treats a published brain encoder as a measurement device whose inverse is to be characterised. Figure 4.1 shows the complete pipeline.

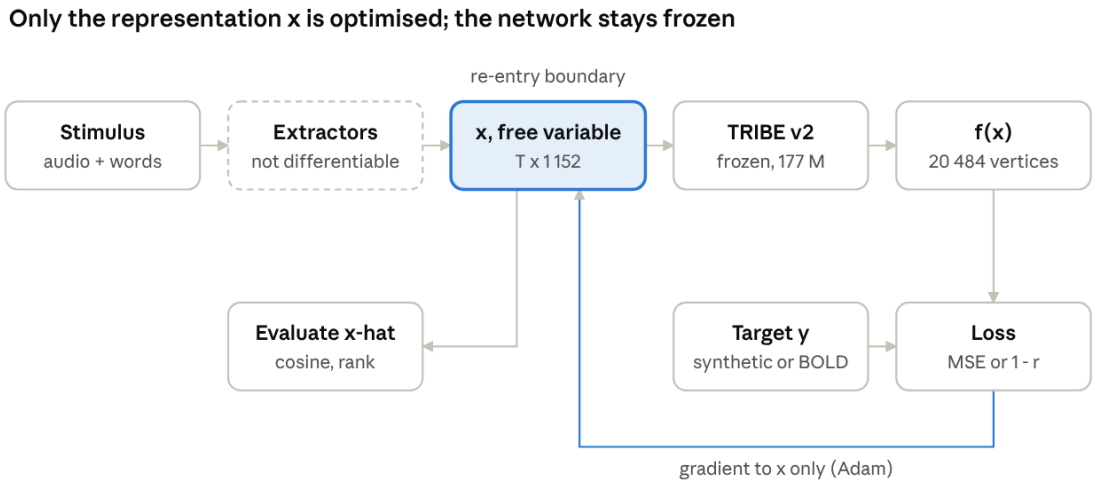


Figure 4.1: Inversion pipeline: the network is re-entered after feature aggregation; only the representation x is optimised, the network stays frozen

4.1 Design Approach / Materials & Methods

4.1.1 The forward operator: TRIBE v2

TRIBE v2 [6] maps a multimodal stimulus representation x ∈ ℝ^(T×1,152), made of text, audio and video slots of 384 dimensions each sampled at 2 Hz, to a predicted response at V = 20,484 vertices of the fsaverage5 cortical surface [20], 100 output steps at 1 Hz. The network has 177.21 M parameters: a transformer encoder (127.47 M), a low-rank head (1,152 → 2,048), a subject-layer predictor (41.97 M) and a temporal pooler. The public release ships a single averaged “unseen subject” layer. All weights are frozen in the proposed work.

4.1.2 Training data of TRIBE v2

Every conclusion depends on the encoder, so the data it learned from is stated here. TRIBE v2 was trained on four public datasets totalling 25 healthy adults and 451.6 hours of 3T fMRI, recorded while participants passively watched or listened to naturalistic stimuli with no task (Table 4.1). The figures come from Table 1 of the TRIBE v2 paper [6] and the study definitions shipped with its code.

Table 4.1: TRIBE v2 training data: participants and stimuli

Dataset	Task in the scanner	Subjects	Sessions	fMRI (h)	TR (s)
CNeuroMod / Algonauts 2025 [10]	Watched the sitcom Friends (seasons 1–6) and four feature films; video, audio and dialogue	4	1,408	268.7	1.49
BOLD Moments [21]	Watched 1,102 three-second naturalistic video clips (1,000 training clips; 102 test clips repeated 10 times)	10	520	61.9	1.75
LeBel et al. 2023 [22]	Listened to spoken autobiographical stories (podcasts); 82 stories for three subjects, 26–27 for the others	8	432	85.8	2.0
Wen et al. 2018 [23]	Watched natural movie clips in 8-minute segments; test segments repeated 10 times	3	258	35.2	2.0
Training total	**All 3T, passive viewing or listening**	**25**	**2,618**	**451.6**	**—**

The model was then evaluated, without further training, on four datasets it had never seen (Table 4.2), bringing the total to 720 subjects, 5,094 sessions and 1,117.7 hours of fMRI [6].

Table 4.2: TRIBE v2 evaluation data (never used for training)

Dataset	Stimuli	Scanner	Subjects	fMRI (h)
NNDb (Aliko et al., 2020)	Full-length films (audio, video, text)	3T	86	160.6
Le Petit Prince (Li et al., 2022)	Audiobook (audio, text)	3T	112	180.2
Narratives (Nastase et al., 2021)	Spoken stories (audio, text)	3T	321	146.6
HCP (Van Essen et al., 2013)	Movie clips (audio, video, text)	7T	176	178.7
Evaluation total			**695**	**666.1**

Two consequences for this project follow. First, the real-data subject (UTS03) is one of the eight LeBel training subjects, but the test story, wheretheressmoke, was excluded from TRIBE v2's training (it is part of the test set of the original TRIBE v2 study), so the encoder never saw the responses that are inverted here; this report calls it the unseen test story. Second, because the release averages subjects, the population model is inverted rather than a subject-specific one.

4.1.3 Choice of fMRI over EEG

EEG is cheaper, portable and has millisecond resolution, so the choice of fMRI needs justification. It rests on five points (Table 4.3). (1) Spatial detail: fMRI resolves millimetre-scale cortical patterns at 20,484 vertices [24], while 32–128 scalp electrodes see a spatially blurred mixture because of volume conduction. (2) A foundation encoder to invert exists only for fMRI; no open whole-brain encoder of naturalistic EEG at comparable scale exists. (3) Decoding through EEG would add a second ill-posed inverse problem, source localisation [25], in front of the one this project studies. (4) Empirically, the same decoder reaches a 32% character error rate from MEG but 67% from EEG [13]. (5) Open naturalistic language data

with fMRI is plentiful (LeBel: 85.8 h), whereas the open simultaneous EEG–fMRI dataset contains rest, an abstract animation and short movie clips but no sustained speech [26]; this ruled out the EEG–fMRI fusion option considered in Report I.

Table 4.3: fMRI versus EEG for this project

Criterion	fMRI (proposed work)	EEG
Spatial detail	≈2–3 mm voxels; 20,484 cortical vertices	32–128 scalp channels; centimetre-scale, blurred
Foundation encoder to invert	TRIBE v2: 177 M parameters, 1,117.7 h, 720 subjects	None at comparable scale
Number of inverse problems	One: response → representation	Two: scalp → sources (ill-posed) → representation
Measured language decoding	Continuous language from fMRI [11]	67% character error vs 32% with MEG [13]
Open naturalistic language data	LeBel 2023: 8 subjects, 85.8 h	No open EEG–fMRI set with sustained speech [26]
Temporal resolution	1.5–2 s samples, ≈5 s haemodynamic lag (cost)	Milliseconds (advantage)
Cost and portability	Scanner-bound (cost)	Wearable, low cost (advantage)

EEG's advantages matter for a deployed device, not for measuring how invertible an encoder is. The method developed here transfers to EEG once an EEG encoder of comparable scale exists.

4.1.4 Inversion method

The public TRIBE v2 inference path is not differentiable end to end: stimuli pass through speech synthesis, transcription and frozen feature extractors, and prediction runs under PyTorch's inference mode. Inspection of the forward pass showed that the input batch is consumed at one point only, the aggregation of extractor features into x; everything after it is tensor arithmetic. Therefore x is taken as the free variable and the network is re-entered at that boundary, f(x) = pool(predictor(head(encoder(x)))). Gradients reach every entry of x ($\|\partial y/\partial x\| = 7.97 \times 10^3$, no zero entries), and re-entering at this boundary reproduces the released model's predictions exactly (r = 1.0000). The representation is recovered by

$$\hat{x} = \operatorname{argmin}_x \|f(x) - y\|^2 + \lambda R(x) \qquad (4.1)$$

using Adam [27] with learning rate 5×10^{-2} and cosine annealing over the step budget, initialisation $0.01 \cdot N(0, I)$, and R an ℓ_2 or temporal-smoothness penalty where tested ($\lambda = 0$ otherwise). On real fMRI, where the scale and offset of the measured response are arbitrary, the loss is one minus the Pearson correlation over time, averaged over observed vertices.

4.2 Codes and Standards

Table 4.4: Standards and frameworks the project follows

Standard / framework	Scope	Alignment of this project
BIDS [28]	Neuroimaging data organisation	Real data from OpenNeuro ds003020, a BIDS dataset
IEEE P2731 [29]	Unified BCI terminology (functional model)	Pipeline described as acquisition → features → decoding → output
IEEE P2794 (draft)	Reporting standards for in-vivo neural interface research	Report states model, data, preprocessing, metrics and failures
FAIR principles [30]	Findable, accessible, interoperable, reusable research outputs	Seeded scripts, logged results, public code and notebooks
OECD Recommendation [31]; UNESCO Recommendation [32]	Responsible innovation and ethics of neurotechnology	Mental-privacy implications discussed; no new human data collected
Licences	TRIBE v2 (CC BY-NC 4.0), Llama 3.2 licence, OpenNeuro data	Non-commercial academic use; gated access obtained

4.3 Constraints, Alternatives and Tradeoffs

Report I compared three solutions: an end-to-end supervised decoder (needs large per-subject training data and gives little insight into why decoding fails), inversion of a pretrained encoder (selected), and a simultaneous EEG–fMRI fusion decoder (blocked by the lack of open simultaneous data with sustained language). The main constraints and the trade-offs taken are listed in Table 4.5.

Table 4.5: Constraints and the trade-offs adopted

Constraint	Effect	Trade-off adopted
No local GPU	Large models cannot run on the team's laptops	Free Colab T4 for GPU work; CPU for regional analysis; Llama in half precision
12 GB RAM / 15 GB GPU on Colab	Llama-3.2-3B and w2v-BERT cannot be loaded together	One story per process; audio and text features in separate passes
Released pipeline not differentiable	Raw stimuli cannot be optimised	Invert the aggregated representation x (one stage after the extractors)
No ground truth in real fMRI	Recovery error cannot be measured directly	Exact synthetic evaluation, plus identification with a known chance level
Only an averaged subject layer released	No per-subject analysis	Invert the population model; state as a limitation
Synthetic x in a random subspace	Absolute values may not transfer	Real-data validation confirms the qualitative pattern

5. PROJECT DESCRIPTION

The project was carried out in five stages, each building on the previous one.

5.1 Instrumentation of the encoder

TRIBE v2 was loaded directly from its configuration and checkpoint, bypassing the text-to-speech and transcription pipeline. The forward map was rebuilt from the aggregation boundary (encoder, low-rank head, predictor, pooler), gradient flow to x was verified, and the rebuilt map was checked against the released predictions on real stimuli (r = 1.0000).

5.2 Synthetic evaluation with exact ground truth

Construction. The ground truth must be a valid input of the operator being inverted, so it lives in TRIBE v2's own representation space: $x \in \mathbb{R}^{(32 \times 1{,}152)}$, 32 steps at 2 Hz (16 s) across the text, audio and video slots, exactly the tensor the feature extractors hand to the encoder. Eight unit directions $b_1 \ldots b_8$ are drawn at random in $\mathbb{R}^{1{,}152}$ (random orientation, so no direction is privileged), and each carries its own cosine time course, $x(t) = \Sigma_k \cos(2\pi f_k t)\, b_k$ with f_k = 2, 4, …, 16 cycles per window (0.125–1 Hz); the result is scaled to unit standard deviation, the scale of real features. Four seeds (0–3) give four independent ground truths, and white noise (i.i.d. Gaussian, full rank) serves as an off-manifold control. The target is y = f(x), so recovery is measured exactly: cosine to x per time step, relative error, encoding error and retrieval rank among time steps. Directional sensitivity along a unit direction v is a finite-difference probe of the Jacobian:

$$s(v) = \|f(x + \varepsilon v) - f(x)\| / \varepsilon \qquad (5.1)$$

5.2.1 Justification and validity of the ground truth

Correct by construction. A known-answer test is the standard way to evaluate an inverse method [8], [17]: because x is chosen and y = f(x) is produced by the frozen model itself (whose re-entry reproduces the released predictions, r = 1.0000), every recovery number is an exact measurement of the operator's inverse, with no modelling assumption about the brain. The only point needing justification is whether the chosen x resembles real stimuli closely enough for the conclusions to transfer. Four pieces of evidence address that (Table 5.1).

(1) Real representations are low-dimensional, like the synthetic truth. The aggregated representations of the six LeBel stories (audio slot, 246 windows of 16 s) have an effective dimensionality (participation ratio) of 4.9 ± 2.1, with 84 ± 6% of their variance in their top

eight directions and a standard deviation of 1.10; the synthetic truth has 7.8, 100% and 1.00, and white noise 30.2 and 31% (Figure 5.1). An 8-D subspace is therefore the right order of magnitude, and slightly harder than real features, not easier. **(2) The conclusions do not depend on its shape.** White noise recovers as well as the structured truth (0.947 vs 0.942), and changing the rate of change 64-fold, which spans the real lag-1 autocorrelation of 0.15 ± 0.12, leaves recovery at 0.836–0.858 with no trend. **(3) They replicate across four independent ground truths** (agreement minimum 0.894 ± 0.003, barrier 4.8×, sensitivity 2.15 ± 0.17×).

(4) A real representation as ground truth gives the same picture. The tests were repeated with a real 16-s window of the unseen test story (whereatheressmoke, 16–32 s) as x, inverted from y = f(x) through the frozen model from three blind starts and one near-truth control (3,000 steps on CPU, scripts/real_x_groundtruth.py). Recovery was partial and start-dependent (cosine 0.43–0.60 in the filled audio slot, against about 0.56 for the synthetic truth at the same budget); the blind solutions disagreed (pairwise cosine 0.19–0.56) and were separated by barriers (the error rose 11–21× on the straight lines between them); and the near-truth control found a deeper minimum (2.9× lower error, recovery 0.95) 54× barrier away (Figure 5.2). The sensitivity probe was inconclusive at this budget because the blind fits had not yet converged (their errors still differed 2.2-fold), so it is reported only for the converged synthetic runs. Real stimulus representations thus show the same non-convex, start-dependent inverse, more strongly; the synthetic truth is the conservative case.

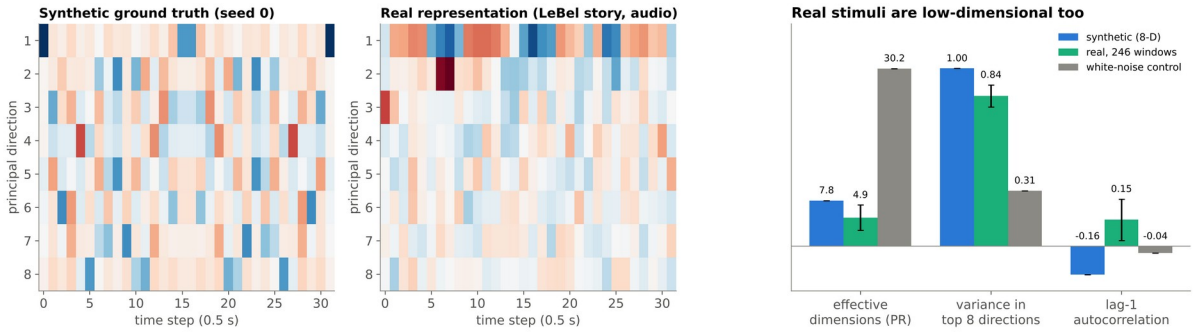


Figure 5.1: Synthetic ground truth and a real stimulus representation along their top eight principal directions, and their effective dimensionality, variance in eight directions and lag-1 autocorrelation (real: 246 windows of the six LeBel stories)

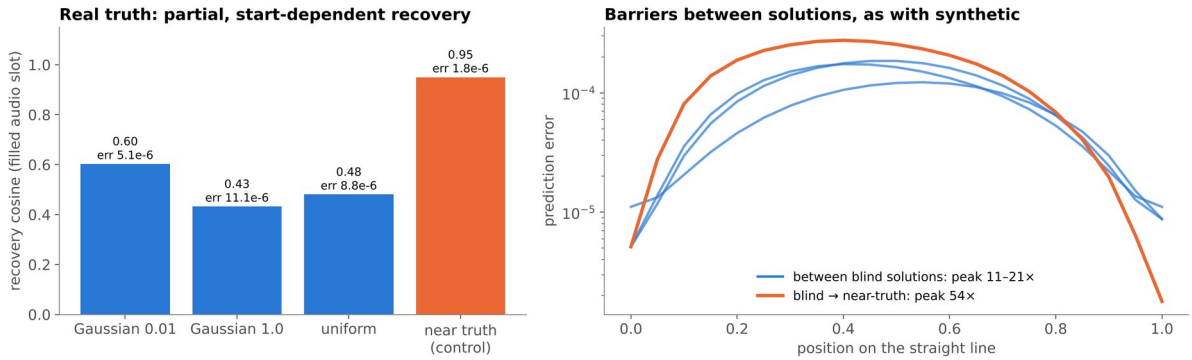


Figure 5.2: Inversion with a real stimulus representation as ground truth: recovery per start (left) and prediction error along the straight lines between solutions (right)

Table 5.1: Evidence that the synthetic ground truth is valid

Check	Test	Result
Exactness of the recovery error	Ground truth chosen; y = f(x) from the frozen model; re-entry reproduces the release	Exact by construction (r = 1.0000)
Realism of the 8-D subspace	Effective dimensionality of real story representations	Real 4.9 ± 2.1 (84% in 8 directions) vs synthetic 7.8; white noise 30.2
Realism of the scale	Standard deviation of real features	1.10 vs 1.00
Effect of structure	White-noise (full-rank) ground truth	0.947 vs 0.942: no effect
Effect of the speed of change	Rate of change varied 64-fold (spans real lag-1 0.15 ± 0.12)	0.836–0.858, no trend: no effect
Dependence on a single ground truth	Four independent seeds	Agreement 0.894 ± 0.003; barrier 4.8×; sensitivity 2.15 ± 0.17×
Behaviour with a real stimulus	Real 16-s window as ground truth, 3 blind starts + control	Partial recovery 0.43–0.60; disagreement 0.19–0.56; barriers 11–21×; deeper minimum 54× away

5.3 Discriminating the causes of incomplete recovery

Four explanations were tested, each with a signature that would appear if it were true. (H1) Information loss: fit-matched solutions differ along a null-space direction the operator ignores. (H2) Non-convexity: distinct minima at equal fit, separated by a barrier. (H3) Optimiser artefact: recovery keeps rising with budget. (H4) Unidentifiable gain: a rescaled x gives the same output. Inversions were started from eight initialisation distributions and compared only when within 2× of the best fit, so under-converged runs could not pose as alternative solutions; the analysis was replicated over four ground truths. For two fit-matched solutions, a null space predicts that the prediction error stays flat along the straight line between them, while separate minima predict a rise in between; this test and the sensitivity probe use forward passes only and do not depend on the optimiser. Two further starts placed near the truth serve as a control for whether a deeper minimum exists.

5.4 Regional analysis

Six disjoint networks were defined on fsaverage5 from the Destrieux atlas [33] (Figure 8.10): left language network [34] (1,282 vertices), its right-hemisphere homologue (1,368), auditory cortex (277), default-mode network [35] (1,224), visual cortex (1,582) and somatomotor cortex (1,487) as a low-semantic control. For each network, the analysis measured forward sensitivity per input modality against 2,000 random vertex sets of equal size, recovery when only that network is observed (8,000 steps) against size-matched random vertex sets with a second random seed as the control spread, and the participation ratio [36] of its responses to 128 random input directions.

5.5 Real-fMRI validation

Data. Subject UTS03 of the LeBel et al. corpus [22] (OpenNeuro ds003020): six stories, 66.6 minutes of speech, 10,363 words, TR 2 s. The unseen test story, wheretheressmoke, is the mean of ten repeats; the other five stories are used for every tuning choice.

Surface projection. The authors' preprocessed responses (95,556 cortical voxels) were mapped to fsaverage5 with TRIBE v2's own projector settings: a 3-mm ball of 20 points around each mid-cortical vertex, trilinear weights inside the cortical mask. Vertices were carried from MNI space to the subject's T1 by an ANTs SyN registration [37] and into functional voxels by the pycortex transform [38], giving a sparse 20,484 × 95,556 sampling matrix; the series were then detrended and z-scored. Validation: the left–right homotopic correlation was 0.635 against −0.017 for heterotopic pairs (above 97% of alternative pairings), and a scrambled-vertex control fell to chance.

Features and alignment. Stimulus features came from TRIBE v2's own pipeline with the dataset's word timings: w2v-BERT 2.0 for audio [39] and Llama-3.2-3B for text [40], grouped into 100-s segments (45 across six stories). The offset between model time and BOLD was the lag maximising prediction–response correlation on the five training stories; the 2,000 analysis vertices were those best predicted on the training stories.

Evaluation. (i) Encoding accuracy per vertex on the unseen test story, against a split-half noise ceiling from the ten repeats. (ii) Forward identification: each 20-s window of the test-story BOLD matched to predictions for all 194 windows of the six stories. (iii) Inversion: each 100-s test segment inverted from measured BOLD, and $\hat{x}$ ranked against the true representations of all 45 segments (chance 1/45). (iv) An oracle that inverts the model's own noise-free prediction under the same budget, separating the operator's conditioning from

measurement noise.

6. HARDWARE/SOFTWARE TOOLS USED

Table 6.1: Hardware and software tools, their use and their cost to the project

Category	Tool	Use in this project	Cost (INR)
Compute (GPU)	Google Colab free tier, NVIDIA T4 (15 GB), 12 GB RAM	Synthetic characterisation; real-data text and audio features and inversion	0 (free tier)
Compute (CPU)	Laptop and cloud CPU (Linux)	Regional analysis, surface projection, audio-only real-data run, figures	0 (existing)
Programming	Python 3, Jupyter notebooks	All analysis code and the Colab notebooks	0 (open source)
Deep learning	PyTorch [41], Hugging Face Transformers and Hub	Autograd through the frozen encoder; Llama-3.2-3B and w2v-BERT	0 (open source)
Encoder	TRIBE v2 code and weights (Meta FAIR)	Frozen forward operator	0 (open weights, CC BY-NC 4.0)
Neuroimaging	nilearn [42], nibabel, ANTsPy, pycortex	Atlas, surfaces, registration, surface projection	0 (open source)
Scientific Python	NumPy, SciPy, scikit-learn, h5py, pandas, Matplotlib	Analysis, ridge baseline, data I/O, figures	0 (open source)
Data	OpenNeuro ds003020 (LeBel et al. 2023)	Real naturalistic fMRI	0 (open data)
Version control and release	Git, GitHub	Code release, version history, public repository	0 (free plan)
Storage	Google Drive	Model weights, features and results for Colab	0 (free quota)
Documentation	Microsoft Word, Microsoft PowerPoint (Microsoft 365 through the VIT student account)	Report, review slides and A3 poster	0 (institutional licence)
Document conversion	LibreOffice	PDF export and page checks	0 (open source)
Total software cost			**0**

All software used is open source, free of charge at the tier used, or covered by the institutional licence, so the software cost of the project is zero; the only expense is printing (Chapter 11).

No hardware was built: the project is computational. All code, results and notebooks are public at github.com/Ivxn404/semantic-decoding-tribe-inversion (Appendix A).

7. SCHEDULE AND MILESTONES

Table 7.1: Schedule and milestones

Stage	Activity	Period (2026)	Outcome
S1	Literature review and problem formulation	Aug	Scoped problem statement (Review I)
S2	Model instrumentation; inversion boundary	Aug – Sep	Working differentiable inverse
S3	Synthetic characterisation (budget, noise, extent, priors)	Sep	Dissociation and conditioning results
S4	Controls and hypothesis tests	Sep – Oct	Non-convexity shown; null space rejected (Review II)
S5	Regional analysis on six networks	Oct	Predictable-vs-recoverable dissociation
S6	Real-fMRI validation (LeBel 2023)	Oct	Identification on an unseen test story (Review III)
S7	Manuscript, report, poster, code release	Oct	IEEE-format manuscript; public repository; final review

Reviews: Review I (problem and plan), Review II on 14 October 2026 (40% milestone), Review III (60–70% milestone), and the final review with poster, viva and demonstration on 29 October 2026.

8. RESULT ANALYSIS

8.1 Encoding accuracy does not certify decoding

Independent inversions at budgets from 250 to 32,000 steps show encoding error falling 28.4-fold while recovery saturates (Figure 8.1, Table 8.1). The final halving of encoding error (5.65 × 10^{-3} → 2.66 × 10^{-3}) buys only +0.015 in recovery cosine. Forward accuracy and decodability are separate quantities.

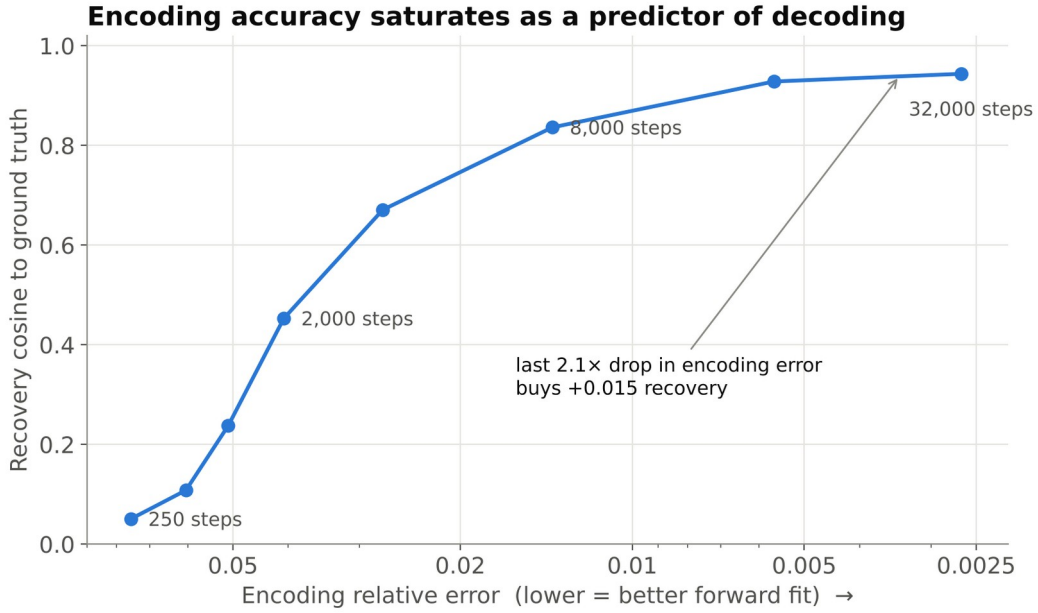


Figure 8.1: Encoding error against recovery cosine across independent step budgets

Table 8.1: Encoding error and recovery across optimisation budgets

Steps	Encoding relative error	Recovery cosine
250	7.546 × 10^{-2}	0.0505
500	6.043 × 10^{-2}	0.1081
1,000	5.098 × 10^{-2}	0.2374
2,000	4.073 × 10^{-2}	0.4529
4,000	2.733 × 10^{-2}	0.6708
8,000	1.380 × 10^{-2}	0.8362
16,000	5.649 × 10^{-3}	0.9282
32,000	2.658 × 10^{-3}	0.9436

8.2 Equally good fits give different answers

At 16,000 steps, seven of eight initialisations reached matched fit (2.52–2.96 × 10^{-7}; one was excluded at 3.4× the best fit). Their mean pairwise cosine was 0.927, minimum 0.835, and mean cosine to the truth 0.914. The effect replicated across four ground truths with the same worst pair each time (Figure 8.2, Table 8.2); recovered norm ratios clustered by initialisation

to within about 0.01, so distinct solutions are reached systematically, not by chance.

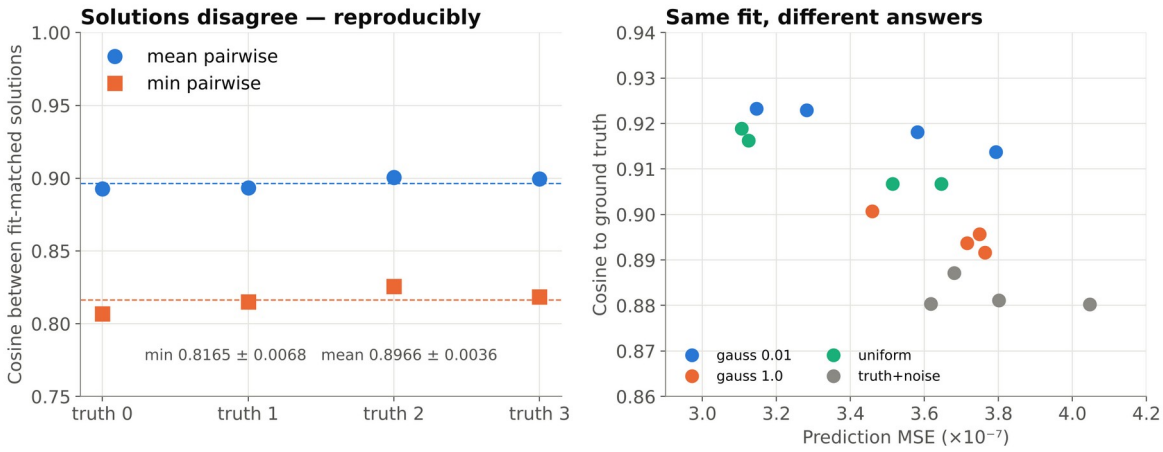


Figure 8.2: Agreement among fit-matched solutions across four ground truths (left) and fit versus recovery per initialisation (right)

Table 8.2: Agreement among fit-matched solutions, four ground truths, 14,000 steps

Ground truth	Best fit (MSE)	Mean pairwise cosine	Min pairwise cosine
0	3.51×10^{-7}	0.893	0.807
1	3.65×10^{-7}	0.893	0.815
2	3.11×10^{-7}	0.901	0.826
3	3.13×10^{-7}	0.900	0.819
Mean ± SD		**0.897 ± 0.004**	**0.817 ± 0.007**

8.3 Barriers, not a null space

The straight line between two fit-matched solutions distinguishes the hypotheses. If the solutions lay in one flat valley (a null space), the prediction error would stay at the endpoints' level all along the line; if they are separate minima, the error rises in between. The test uses forward passes only. It was run on 19–20 October 2026 with notebooks/Verification_Tests.ipynb (14,000 steps, four ground truths), comparing only the five starts that were blind to the truth and fit-matched (within 1.3× of the best blind fit in every case; agreement mean 0.955, minimum 0.894 ± 0.003 over the four truths).

Over the 40 pairs, the error rose between every pair of distinct solutions (29 pairs with agreement below 0.99) to 2.4–8.2× the worse endpoint (median 4.8×); the 11 near-identical pairs (agreement ≥ 0.99, mostly the zero and small-Gaussian starts) lie in the same minimum and stay nearly flat (1.0–2.2×). The rise grows as the two solutions become less similar (correlation −0.70; Figure 8.3, Table 8.3). Two control starts placed near the truth reached a 4.2 ± 0.3× lower error and higher recovery (0.946–0.977, against 0.891–0.926 for the blind starts), and the straight path from the best blind solution to the near-truth one rose 8.1 ± 3.2× (4.7–12.4×) above the blind end. The blind solutions are therefore distinct local minima,

separated by barriers from one another and from a deeper minimum near the truth: the inverse is non-convex, and the information-loss hypothesis is rejected.

Correction. An earlier version of this test chose the least similar pair among all eight starts and compared the change along d with random directions. That rule paired an under-converged solution (Gaussian s.d. 5, error 1.15 × 10^{-6}, recovery 0.62), so its numbers (2.6× falling to 0.6×) are withdrawn. With a fit-matched pair, the change along d stays far below random directions at every scale (2.6×, 6.0×, 13×, 258× at 0.1, 0.25, 0.5 and 1.0 of the gap), as it must when both ends fit the target; that comparison cannot detect a barrier, which is why the straight-path error is reported instead. The selection rule in scripts/synthetic_characterisation.py has been corrected.

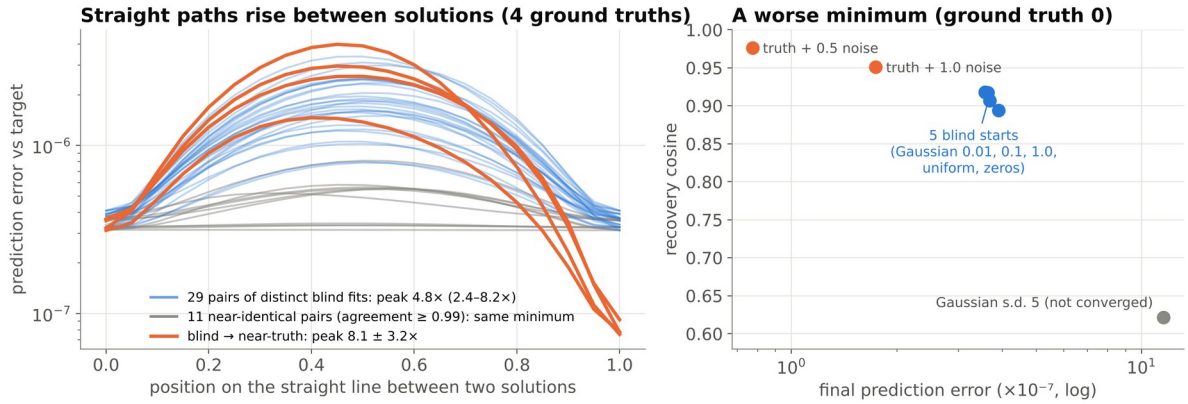


Figure 8.3: Left: prediction error along the straight line between fit-matched blind solutions (blue, ten pairs) and from the best blind solution to the near-truth solution (orange). Right: final error and recovery for each of the eight starts

Table 8.3: Straight-path barrier between fit-matched blind solutions: all ten pairs for ground truth 0, and the summary over four ground truths (14,000 steps)

Solution pair	Agreement (cosine)	Peak error on the path	Peak / worse end
Gaussian 1.0 – zeros	0.927	2.36 × 10^{-6}	6.0×
Gaussian 1.0 – uniform	0.890	1.79 × 10^{-6}	4.6×
Gaussian 0.01 – Gaussian 1.0	0.927	1.71 × 10^{-6}	4.4×
uniform – zeros	0.958	1.50 × 10^{-6}	4.1×
Gaussian 0.1 – Gaussian 1.0	0.940	1.57 × 10^{-6}	4.0×
Gaussian 0.01 – uniform	0.957	1.32 × 10^{-6}	3.6×
Gaussian 0.1 – uniform	0.956	1.25 × 10^{-6}	3.4×
Gaussian 0.01 – Gaussian 0.1	0.992	8.12 × 10^{-7}	2.2×
Gaussian 0.1 – zeros	0.993	5.52 × 10^{-7}	1.6×
Gaussian 0.01 – zeros	0.999	5.57 × 10^{-7}	1.5×
Best blind (zeros) – near-truth	—	2.58 × 10^{-6}	7.2×
Four truths: 29 distinct pairs (agreement < 0.99)	0.89–0.99	—	median 4.8× (2.4–8.2×)
Four truths: 11 near-identical pairs	≥ 0.99	—	1.0–2.2×
Four truths: best blind – near-truth	—	—	8.1 ± 3.2×

8.4 Errors lie in the operator's weakest directions

On the least similar fit-matched pair (Gaussian 1.0 and uniform), sensitivity along their difference d was 8.10×10^{-2} for ground truth 0, against a median of 1.68×10^{-1} over 200 random directions (range 1.22–2.39 × 10^{-1}): 2.1× below the median and below every random direction sampled. Over four ground truths the ratio was 2.15 ± 0.17× (1.99–2.40×), and along the actual error $\hat{x} - x$ 1.68 ± 0.17×. Nine of the ten blind pairs lay below the median (1.1–2.1×); the exception is the pair of almost identical solutions from the zero and 0.01 starts (agreement 0.999), whose difference is too small to probe reliably. The original run (Figure 8.4) gave 1.9×. The inverse is anisotropically ill-conditioned: the operator responds about half as strongly to the components that inversion fails to recover.

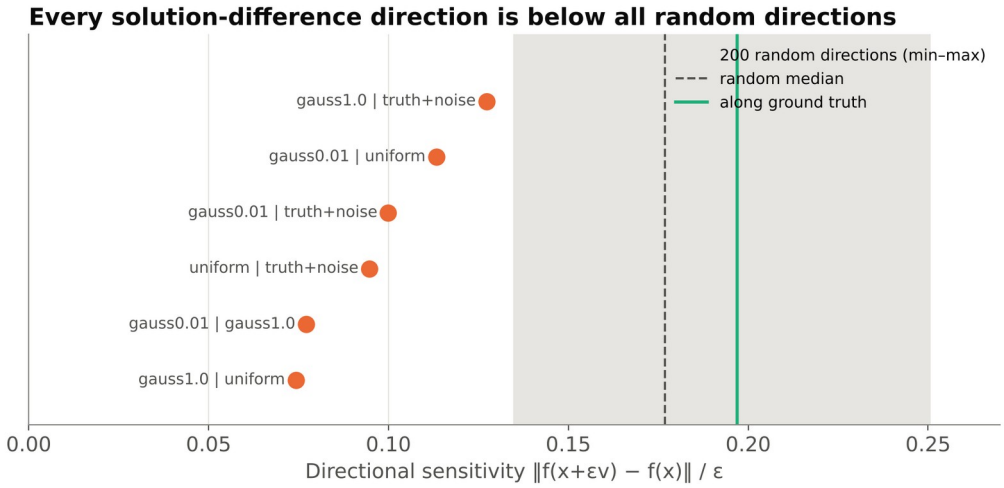


Figure 8.4: Directional sensitivity of the operator in the original run: solution-difference directions against 200 random directions (the corrected pairs give 2.15 ± 0.17× over four truths)

8.5 Further conditioning results

Gain is identifiable: only the true scale lies within the optimiser's residual floor, and a 10% rescale already produces 7.5× the floor (Figure 8.5). The inverse degrades smoothly with measurement noise while retrieval stays near 0.99 (Figure 8.6), and about 10% of cortex (2,000 random vertices) carries most of the recoverable signal (Figure 8.7). Generic priors do not help: an ℓ_2 penalty gives at most +0.032 and temporal smoothness destroys recovery (Figure 8.8). A ridge decoder fitted on the same forward model reaches only 0.060, against 0.918 for gradient inversion (Figure 8.9). Table 8.4 collects these results.

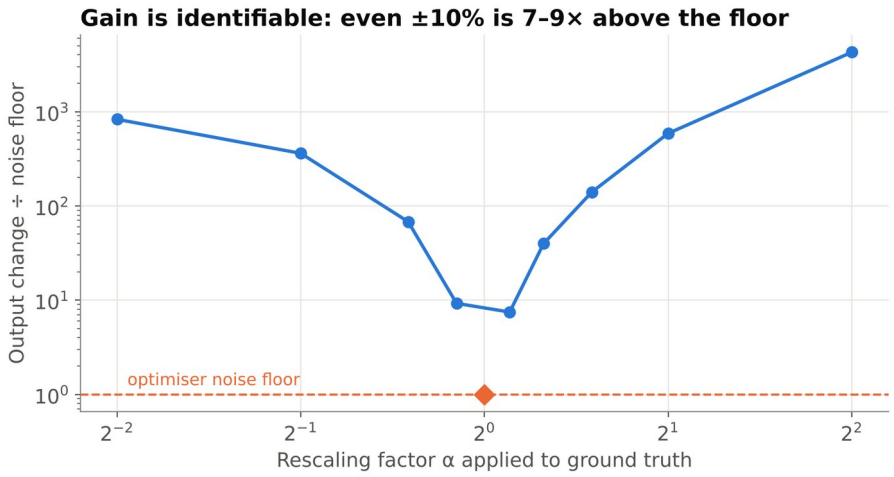


Figure 8.5: Output change under rescaling of the ground truth relative to the optimiser's noise floor

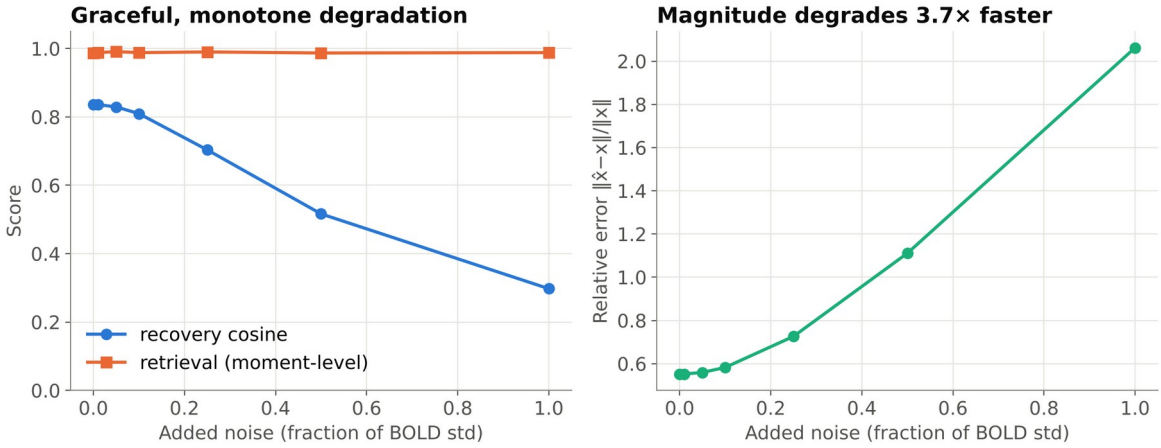


Figure 8.6: Recovery, relative error and retrieval against measurement noise

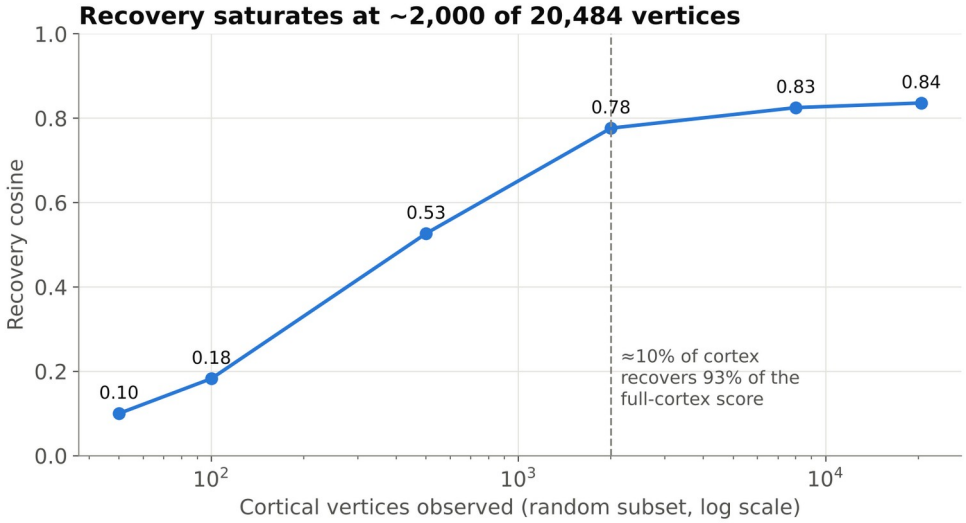


Figure 8.7: Recovery against the number of observed cortical vertices

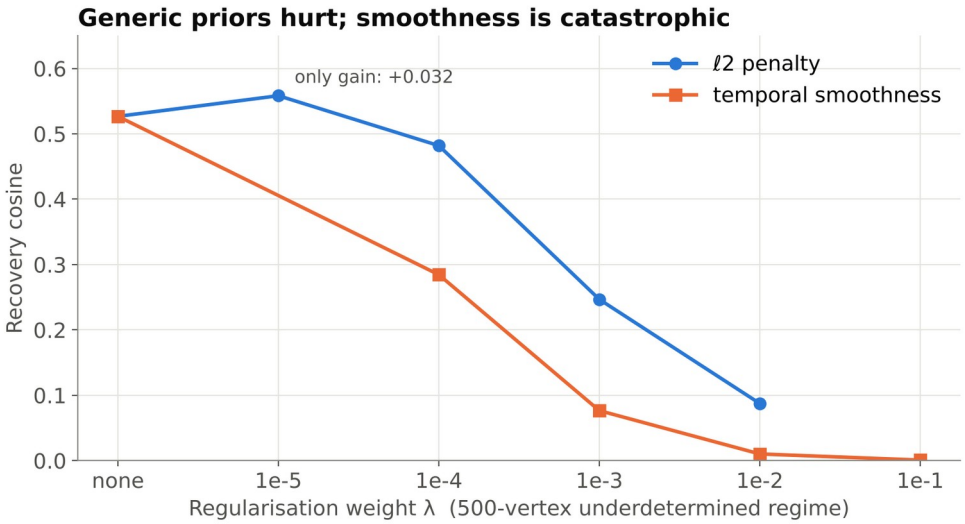


Figure 8.8: Effect of ℓ_2 and temporal-smoothness priors in the under-determined regime (500 vertices)

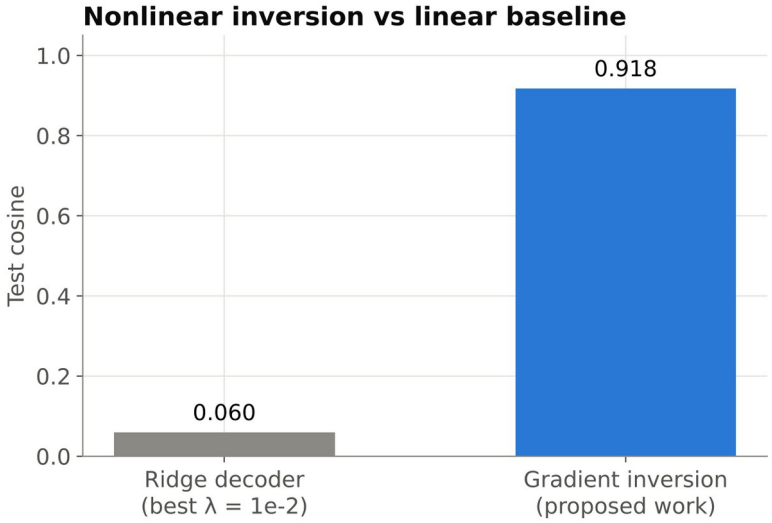


Figure 8.9: Gradient inversion against the best ridge decoder

Table 8.4: Further characterisation of the inverse (whole cortex unless stated)

Test	Result
Gain (rescale x by α)	Only α = 1 within the residual floor; a 10% rescale gives 7.5× the floor
Measurement noise, 0 → 100% of BOLD SD	Recovery 0.836 → 0.298, monotone; retrieval stays ≈0.99
Observed cortex, 50 → 20,484 vertices	Recovery 0.100 → 0.836; 2,000 random vertices (10%) reach 0.777
ℓ_2 prior (500 vertices)	Best +0.032 at $\lambda = 10^{-5}$, harmful above
Temporal-smoothness prior (500 vertices)	Recovery 0.527 → 0.0008; retrieval to chance
Ridge decoder (240 train, 60 test pairs)	Best test cosine 0.060 vs 0.918 for gradient inversion (time-averaged target: a lower bound)
Stimulus manifold	White noise 0.947 ± 0.004 vs smooth 0.942 ± 0.005
Rate of semantic change, 64-fold range	0.836–0.858, no trend (no haemodynamic operator in the release)
Optimiser control (SGD, L-BFGS)	Did not converge (1.4×10^{-4}, 2.8×10^{-5} vs 3.5×10^{-7}); reported as a limitation

8.6 Predictable is not recoverable, region by region

Forward sensitivity follows anatomy (Figure 8.11): language and default-mode networks are dominated by text-derived dimensions (z = +26.6 to +28.5), auditory cortex by audio (+34.7) and visual cortex by video (+40.7). Recoverability does not follow sensitivity (Figure 8.12, Table 8.5). Observing only its own vertices, visual cortex, the most responsive network (1.60× the cortical mean), recovered 0.125 worse than size-matched random vertices; somatomotor cortex, the least responsive (0.73×), recovered 0.102 better and exceeded the whole cortex at the same budget (0.851 vs 0.836). Both effects are about ten times the control spread (second random seed: 0.0014 and 0.0126). Response diversity explains the extremes: somatomotor alone is more diverse than its control (participation ratio 18.6 vs 15.2), while visual cortex is concentrated (10.3; top direction 27% of variance; Figure 8.14).

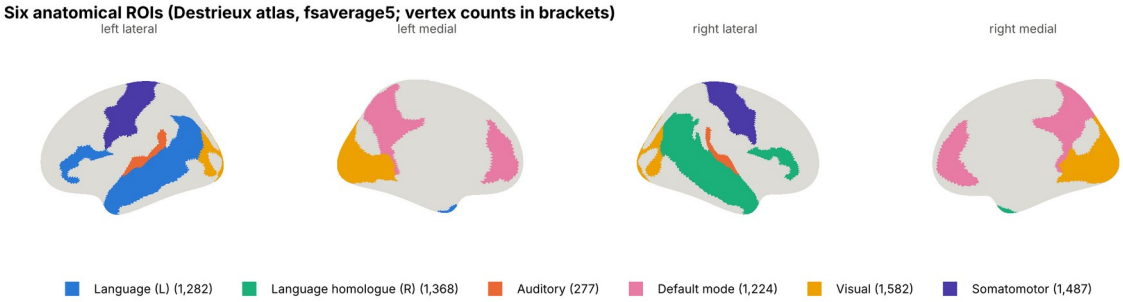


Figure 8.10: The six anatomical networks on the inflated fsaverage5 surface

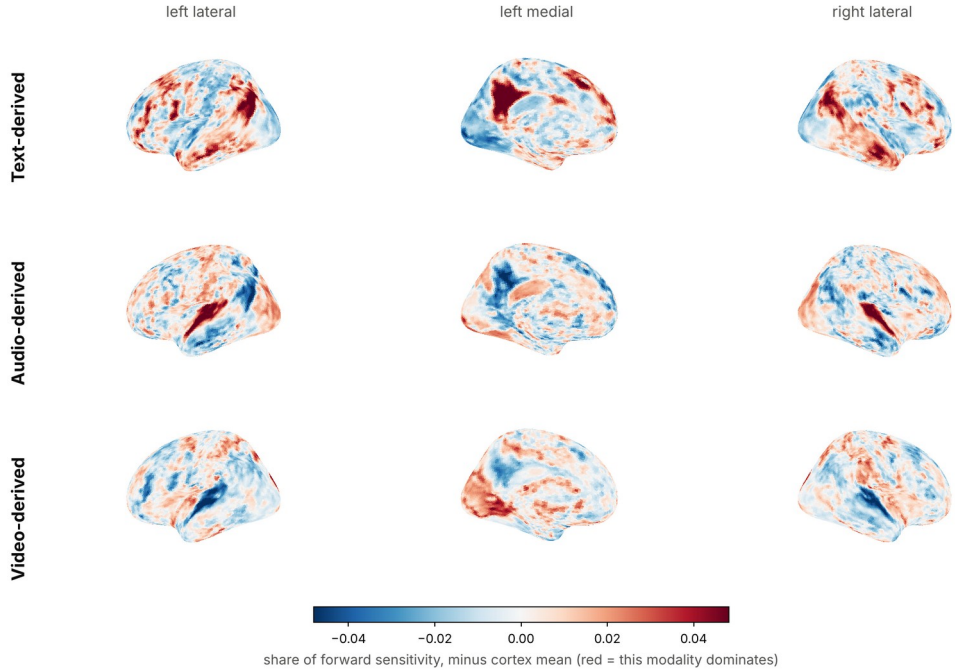


Figure 8.11: Forward-sensitivity share of the text-, audio- and video-derived dimensions across cortex

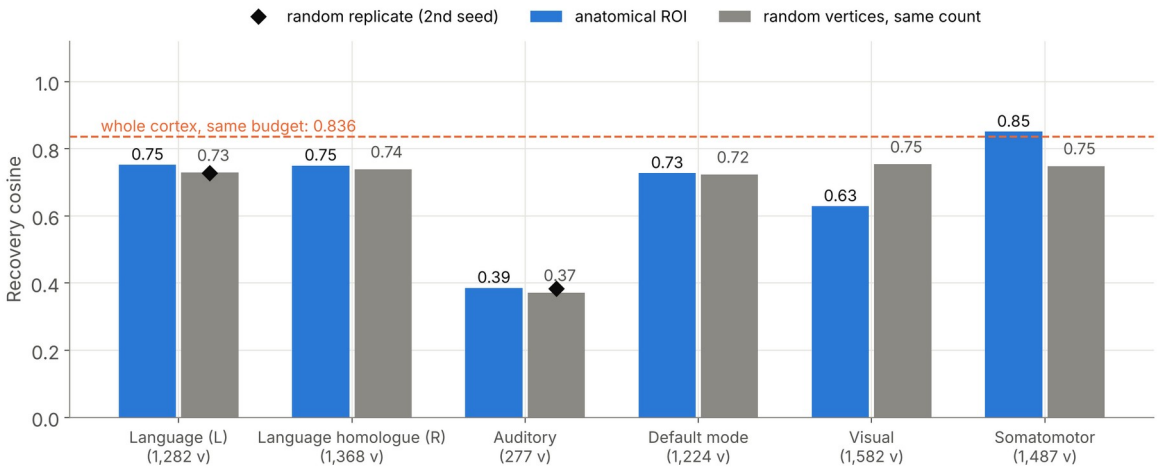


Figure 8.12: Recovery from each network versus size-matched random vertex sets

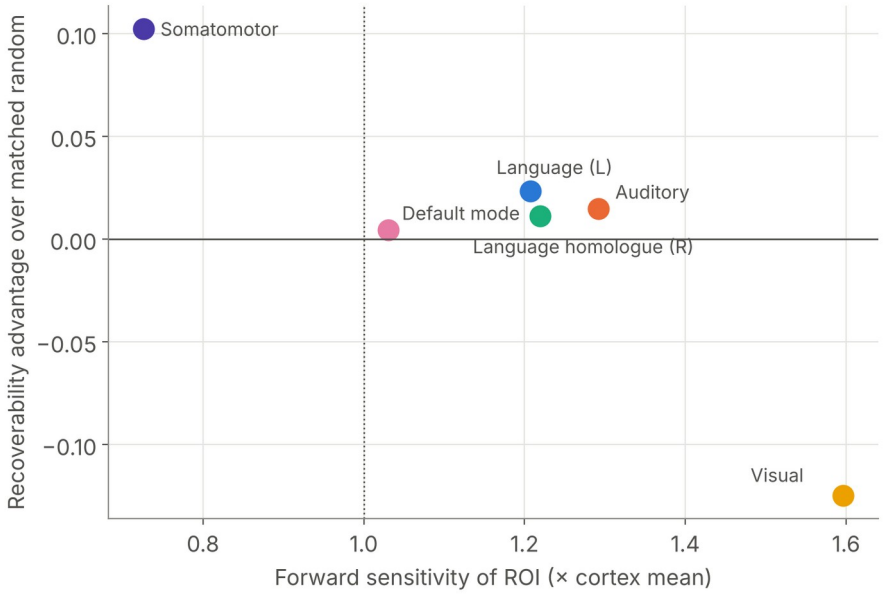


Figure 8.13: Forward sensitivity against recoverability across networks

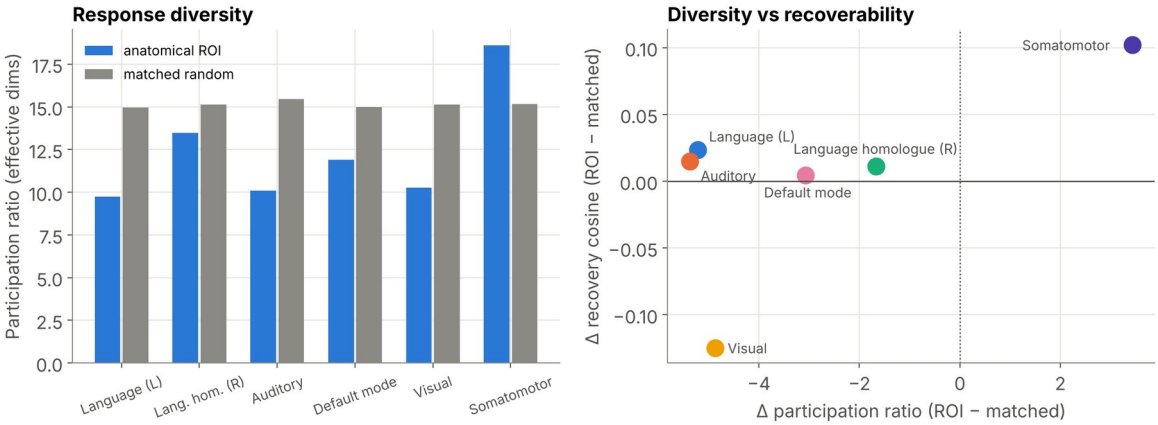


Figure 8.14: Participation ratio of each network's responses against size-matched controls

Table 8.5: Network sensitivity and recoverability (synthetic, 8,000 steps)

Network	Vertices	Sensitivity (× cortex)	Dominant slot (z)	Recovery: network	Recovery: random	Difference	PR (network / control)
Somatomotor	1,487	0.73	none	0.851	0.748	+0.102	18.6 / 15.2
Language (L)	1,282	1.21	text +26.6	0.752	0.729	+0.023	9.7 / 15.0
Auditory	277	1.29	audio +34.7	0.386	0.371	+0.015	10.1 / 15.5
Language (R)	1,368	1.22	text +28.5	0.750	0.739	+0.011	13.5 / 15.1
Default mode	1,224	1.03	text +26.9	0.727	0.723	+0.004	11.9 / 15.0
Visual	1,582	1.60	video +40.7	0.629	0.754	−0.125	10.3 / 15.1

8.7 Real fMRI

The alignment offset fitted on the training stories was 4 s with text and audio features. On the unseen test story the encoder reached r = 0.478 on the 2,000 vertices chosen from the other stories, 56% of the split-half noise ceiling (0.859), led by auditory (0.405), right-language (0.348) and default-mode (0.348) networks (Figure 8.15, Table 8.6). Adding language features more than doubled encoding accuracy compared with audio alone and nearly tripled forward identification of 20-s windows (58.6% against a chance level of 0.5%).

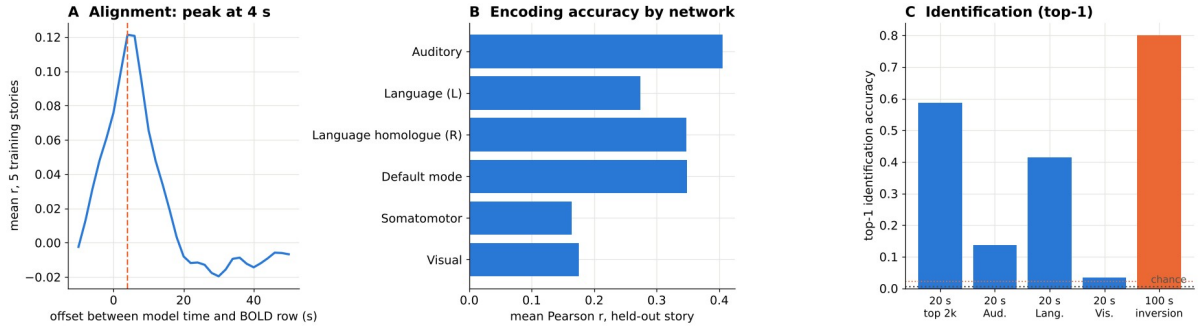


Figure 8.15: Real fMRI, unseen test story: alignment (A), encoding accuracy by network (B) and identification accuracy (C)

Table 8.6: Real-fMRI validation on the unseen test story wheretheressmoke (UTS03)

Measure	Text + audio	Audio only	Chance / ceiling
Encoding r, whole cortex	0.199	0.062	0
Encoding r, top 2,000 vertices	0.478 (56% of ceiling)	0.273 (32%)	ceiling 0.859
Encoding r, auditory / language L / language R	0.405 / 0.274 / 0.348	0.263 / 0.193 / 0.200	0
Encoding r, default mode / somatomotor / visual	0.348 / 0.164 / 0.176	0.103 / 0.039 / −0.020	0
Forward identification, 20-s windows, top-1 (top-5)	58.6% (86.2%)	20.7% (41.4%)	0.5% (194 windows)
Inversion of measured BOLD, 100-s segments, top-1	4 / 5 (ranks 0, 0, 6, 0, 0)	5 / 5	2.2% (45 segments)
Cosine of $\hat{x}$ to the true x	0.024	0.062	—
Oracle: inverting the model's own prediction	5 / 5, cosine 0.224	—	—

Inverting measured BOLD ranked the true segment first in 4 of 5 test queries (the fifth ranked seventh of 45), including against the six other segments of the same story, while $\hat{x}$ kept a cosine of only 0.02 to the truth. The oracle separates two limits: without noise, inversion identifies every segment yet recovers only 0.22 of the representation, so most of the loss is the operator's conditioning, not measurement noise; noise then lowers recovery further while leaving identification largely intact, as in the synthetic noise sweep. On real data, as in simulation, the inverse identifies which moment was heard but does not reconstruct its representation.

8.8 Discussion

Identification is robust; reconstruction is not. Identification needs only the high-sensitivity components that separate one moment from another, and those are the components the operator transmits well; the components it transmits weakly are where recovery fails (Section 8.4). Decoding claims built on identification accuracy should therefore not be read as reconstruction of the underlying representation.

Better encoders need not decode better. A better fit can come almost entirely from high-sensitivity directions, leaving low-sensitivity ones as ill-determined as before. This is the dissociation of Section 8.1, and it means scaling laws for encoding [5] cannot be assumed to carry over to decoding.

The limit is geometry, not lost information. Fit-matched solutions are separate minima, not points on a flat valley, and magnitude is identifiable. Recovery is limited by a non-convex, anisotropically conditioned landscape, which better optimisation, initialisation or priors matched to the operator's weak directions could improve; generic priors do not. For linear inverted encoding models the stimulus is recovered only up to an invertible transform [16]; a deep encoder replaces that closed-form ambiguity with a non-convex, direction-dependent one that must be measured, as done here.

Limitations. Synthetic ground truth lies in a random low-dimensional subspace; inversion targets the aggregated representation; all solutions use Adam; the release has no explicit haemodynamic operator and only the averaged subject layer; regional z-scores use a random-vertex null that ignores spatial autocorrelation; and the real-data test covers one subject and five test inversion queries. A per-network test on real data needs unseen test stories from more subjects, because the other five stories were in TRIBE v2's training set.

9. CONCLUSION

A near-perfect deep brain encoder does not by itself make meaning recoverable from fMRI. Inverting TRIBE v2 is a non-convex, direction-dependently conditioned problem: it identifies which stimulus was presented, on synthetic and real data alike, but reconstructs only the components the operator transmits well. Recoverability has to be measured, not inferred from predictivity, and it differs across cortex in ways predictivity does not reveal.

9.1 Obtained Results

- A differentiable re-entry point into TRIBE v2 that reproduces its predictions exactly ($r = 1.0000$).
- Encoding error falls 28.4× while recovery saturates at cosine 0.944: encoding and decoding dissociate.
- Fit-matched solutions disagree (minimum pairwise cosine 0.817 ± 0.007 over four ground truths): the inverse is non-convex.
- Over four ground truths the prediction error rises 2.4–8.2× (median 4.8×) between every pair of distinct fit-matched blind solutions, and 8.1 ± 3.2× towards a deeper minimum near the truth: barriers, not a null space.
- Solution differences lie 2.15 ± 0.17× below the median sensitivity of 200 random directions (1.9× in the original run): anisotropic ill-conditioning.
- Gain is identifiable; generic priors hurt; gradient inversion beats a linear decoder about 15-fold.
- Visual cortex, the most responsive network, is the worst to decode from (-0.125 vs random); somatomotor, the least responsive, the best ($+0.102$).
- Real fMRI, unseen test story: encoding $r = 0.48$ (56% of the noise ceiling); inversion identifies 4 of 5 segments among 45 (chance 2.2%) with cosine 0.02; the noise-free oracle recovers 0.22.

9.2 Future improvement/work

- Real-data inversion across all eight LeBel subjects and a per-network real-data test on stories excluded from TRIBE v2's training.
- Optimisers and priors matched to the operator's weak directions; multiple restarts and ensembling over minima.

- Inversion of the raw extractor features (one stage earlier) and an explicit haemodynamic model.
- Transfer of the method to EEG and MEG encoders as they appear, and spin-permutation tests for the regional statistics.

9.3 Individual contribution from team members

Table 9.1: Individual contributions

Member	Contribution
Abhranil Sarkar (23BML0077)	Model instrumentation and the differentiable boundary; inversion engine and optimisation; design and execution of the synthetic and control experiments; real-fMRI pipeline (registration, surface projection, features, evaluation); figures, code release, manuscript and reports
Tanisha Choudhuri (23BML0081)	Functional neuroanatomy and definition of the six cortical networks from the Destrieux atlas; curation of the LeBel dataset; interpretation of the regional and real-data findings; contribution to reports, reviews and the poster

10. SOCIAL AND ENVIRONMENTAL IMPACT

Social impact. The work concerns people who keep comprehension but lose speech, for whom a non-invasive communication channel would be transformative. By measuring what fMRI decoding can and cannot recover, it helps direct effort towards approaches that can work and away from over-promising. It also bears on mental privacy: the finding that a near-perfect encoder identifies stimuli but does not reconstruct their representation is relevant to policy discussions of neural data [31], [32]. No new human data were collected; the public data used were released for research with consent obtained by the original studies.

Environmental impact. The project re-used a pretrained model rather than training a new one, which avoids the large energy cost of training a foundation encoder. All computation ran on a single shared cloud GPU and CPUs for a modest number of GPU-hours, and the public release lets others reproduce the results without repeating the experiments from scratch.

11. COST ANALYSIS

The project is computational and required no hardware purchase. Table 11.1 gives the budget estimated in Report I and the cost actually incurred.

Table 11.1: Estimated and incurred cost

Item	Estimate in Report I (INR)	Incurred (INR)	Remark
Hosted GPU compute (subscription)	5,500	0	Free Colab T4 tier was sufficient
Additional on-demand GPU credits	4,000	0	Memory-saving changes avoided paid tiers
Cloud storage (≈200 GB)	1,500	0	Free Google Drive quota and local storage
Software (Python stack, PyTorch, Colab, GitHub, Microsoft Word and PowerPoint, LibreOffice)	0	0	Open source, free tiers and the VIT institutional licence (Table 6.1)
Neuroimaging datasets	0	0	Open licence (OpenNeuro)
Pretrained encoder weights	0	0	Open-weight, non-commercial licence
Documentation and reprographics	1,500	As per printing	Report binding and A3 poster printing
Contingency (≈15%)	2,000	0	Not required
Total	**14,500**	**Printing only**	

Manpower: 580 planned person-hours (330 and 250 hours for the two members), as set out in Report I.

12. PROJECT OUTCOME PUBLICATION/PATENT

- **Manuscript:** “Predictable Is Not Recoverable: Conditioning and Non-Convexity of Inverting a Foundation Brain Encoding Model”, A. Sarkar*, T. Choudhuri* (joint first authors) and S. Nageswaran; complete in IEEE format and ready for submission to a journal or conference (venue to be decided).
- **Open-source release:** github.com/Ivxn404/semantic-decoding-tribe-inversion, with all scripts, logged results, figures, a demonstration video and two Colab notebooks that reproduce the synthetic and real-data analyses.
- **Patent:** none filed; the contribution is a scientific characterisation and an open method.

13. REFERENCES

[1] T. Naselaris, K. N. Kay, S. Nishimoto, and J. L. Gallant, “Encoding and decoding in fMRI,” NeuroImage, vol. 56, no. 2, pp. 400–410, 2011.

[2] K. N. Kay, T. Naselaris, R. J. Prenger, and J. L. Gallant, “Identifying natural images from human brain activity,” Nature, vol. 452, pp. 352–355, 2008.

[3] T. M. Mitchell et al., “Predicting human brain activity associated with the meanings of nouns,” Science, vol. 320, no. 5880, pp. 1191–1195, 2008.

[4] A. G. Huth, W. A. de Heer, T. L. Griffiths, F. E. Theunissen, and J. L. Gallant, “Natural speech reveals the semantic maps that tile human cerebral cortex,” Nature, vol. 532, pp. 453–458, 2016.

[5] R. Antonello, A. Vaidya, and A. G. Huth, “Scaling laws for language encoding models in fMRI,” in Proc. Advances in Neural Information Processing Systems (NeurIPS), vol. 36, 2023.

[6] S. d’Ascoli, J. Rapin, Y. Benchetrit, T. Brooks, K. Begany, J. Raugel, H. Banville, and J.-R. King, “A foundation model of vision, audition, and language for in-silico neuroscience,” arXiv:2605.04326, 2026.

[7] J. Hadamard, “Sur les problèmes aux dérivées partielles et leur signification physique,” Princeton University Bulletin, vol. 13, pp. 49–52, 1902.

[8] P. C. Hansen, Discrete Inverse Problems: Insight and Algorithms. Philadelphia, PA, USA: SIAM, 2010.

[9] S. d’Ascoli, J. Rapin, Y. Benchetrit, H. Banville, and J.-R. King, “TRIBE: TRImodal Brain Encoder for whole-brain fMRI response prediction,” in Proc. Int. Conf. Learning Representations (ICLR), 2026.

[10] A. T. Gifford et al., “The Algonauts Project 2025 challenge: How the human brain makes sense of multimodal movies,” arXiv:2501.00504, 2025.

[11] J. Tang, A. LeBel, S. Jain, and A. G. Huth, “Semantic reconstruction of continuous language from non-invasive brain recordings,” Nature Neuroscience, vol. 26, pp. 858–866, 2023.

[12] A. Défossez, C. Caucheteux, J. Rapin, O. Kabeli, and J.-R. King, “Decoding speech perception from non-invasive brain recordings,” Nature Machine Intelligence, vol. 5, pp. 1097–1107, 2023.

[13] J. Lévy, M. Zhang, S. Pinet, J. Rapin, H. Banville, S. d’Ascoli, and J.-R. King, “Brain-to-text decoding: A non-invasive approach via typing,” arXiv:2502.17480, 2025.

[14] P. S. Scotti et al., “MindEye2: Shared-subject models enable fMRI-to-image with 1 hour of data,” in Proc. Int. Conf. Machine Learning (ICML), 2024.

[15] G. J. Brouwer and D. J. Heeger, “Decoding and reconstructing color from responses in human visual cortex,” Journal of Neuroscience, vol. 29, no. 44, pp. 13992–14003, 2009.

[16] J. L. Gardner and T. Liu, “Inverted encoding models reconstruct an arbitrary model response, not the stimulus,” eNeuro, vol. 6, no. 2, ENEURO.0363-18.2019, 2019.

[17] A. Mahendran and A. Vedaldi, “Understanding deep image representations by inverting them,” in Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR), 2015, pp. 5188–5196.

[18] Bladon and Bent, “Feature visualization recovers known cortical selectivity from TRIBE v2,” arXiv:2605.13904, 2026.

[19] Y. Benchetrit et al., “Boosting brain-to-image decoding with TRIBE v2 data augmentation,” arXiv:2606.06345, 2026.

[20] B. Fischl, M. I. Sereno, R. B. H. Tootell, and A. M. Dale, “High-resolution intersubject averaging and a coordinate system for the cortical surface,” Human Brain Mapping, vol. 8, no. 4, pp. 272–284, 1999.

[21] B. Lahner et al., “Modeling short visual events through the BOLD Moments video fMRI dataset and metadata,” Nature Communications, vol. 15, art. 6241, 2024.

[22] A. LeBel et al., “A natural language fMRI dataset for voxelwise encoding models,” Scientific Data, vol. 10, art. 555, 2023.

[23] H. Wen, J. Shi, Y. Zhang, K.-H. Lu, J. Cao, and Z. Liu, “Neural encoding and decoding with deep learning for dynamic natural vision,” Cerebral Cortex, vol. 28, no. 12, pp. 4136–4160, 2018.

[24] N. K. Logothetis, “What we can do and what we cannot do with fMRI,” Nature, vol. 453, pp. 869–878, 2008.

[25] C. M. Michel, M. M. Murray, G. Lantz, S. Gonzalez, L. Spinelli, and R. Grave de Peralta, “EEG source imaging,” Clinical Neurophysiology, vol. 115, no. 10, pp. 2195–2222, 2004.

[26] Q. K. Telesford et al., “An open-access dataset of naturalistic viewing using simultaneous EEG-fMRI,” Scientific Data, vol. 10, art. 554, 2023.

[27] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in Proc. Int. Conf. Learning Representations (ICLR), 2015.

[28] K. J. Gorgolewski et al., “The brain imaging data structure, a format for organizing and describing outputs of neuroimaging experiments,” Scientific Data, vol. 3, art. 160044, 2016.

[29] IEEE P2731 Working Group, “A functional model for unifying brain–computer interface terminology,” IEEE Open Journal of Engineering in Medicine and Biology, vol. 2, 2021.

[30] M. D. Wilkinson et al., “The FAIR Guiding Principles for scientific data management and stewardship,” Scientific Data, vol. 3, art. 160018, 2016.

[31] OECD, Recommendation of the Council on Responsible Innovation in Neurotechnology, OECD/LEGAL/0457, 2019.

[32] UNESCO, Recommendation on the Ethics of Neurotechnology, adopted 12 November 2025.

[33] C. Destrieux, B. Fischl, A. Dale, and E. Halgren, “Automatic parcellation of human cortical gyri and sulci using standard anatomical nomenclature,” NeuroImage, vol. 53, no. 1, pp. 1–15, 2010.

[34] E. Fedorenko, A. A. Ivanova, and T. I. Regev, “The language network as a natural kind within the broader landscape of the human brain,” Nature Reviews Neuroscience, vol. 25, pp. 289–312, 2024.

[35] R. L. Buckner, J. R. Andrews-Hanna, and D. L. Schacter, “The brain’s default network: Anatomy, function, and relevance to disease,” Annals of the New York Academy of Sciences, vol. 1124, pp. 1–38, 2008.

[36] P. Gao, E. Trautmann, B. Yu, G. Santhanam, S. Ryu, K. Shenoy, and S. Ganguli, “A theory of multineuronal dimensionality, dynamics and measurement,” bioRxiv 214262, 2017.

[37] B. B. Avants, C. L. Epstein, M. Grossman, and J. C. Gee, “Symmetric diffeomorphic image registration with cross-correlation,” Medical Image Analysis, vol. 12, no. 1, pp. 26–41, 2008.

[38] J. S. Gao, A. G. Huth, M. D. Lescroart, and J. L. Gallant, “Pycortex: An interactive surface

visualizer for fMRI," Frontiers in Neuroinformatics, vol. 9, art. 23, 2015.

[39] Y.-A. Chung et al., "W2v-BERT: Combining contrastive learning and masked language modeling for self-supervised speech pre-training," in Proc. IEEE Automatic Speech Recognition and Understanding Workshop (ASRU), 2021, pp. 244–250.

[40] A. Grattafiori et al., "The Llama 3 herd of models," arXiv:2407.21783, 2024.

[41] A. Paszke et al., "PyTorch: An imperative style, high-performance deep learning library," in Proc. Advances in Neural Information Processing Systems (NeurIPS), vol. 32, 2019.

[42] A. Abraham et al., "Machine learning for neuroimaging with scikit-learn," Frontiers in Neuroinformatics, vol. 8, art. 14, 2014.

APPENDIX A: CODE AND REPRODUCTION

Repository: github.com/Ivxn404/semantic-decoding-tribe-inversion (MIT licence for the code). TRIBE v2 weights and the fMRI data are downloaded from their original sources and are not redistributed.

Table A.1: Map from report results to code

Result	Script (scripts/)	Output
Sections 8.1–8.5 (synthetic)	synthetic_characterisation.py	results/synth_*.json (logged numbers: docs/Methods_and_Results.md)
Section 8.3–8.4 verification	notebooks/Verification_Tests.ipynb, make_verification_fig.py	results/verification/
Section 8.6 (regional)	roi_sensitivity.py, roi_share_stats.py, roi_inversion.py, roi_dimensionality.py	results/roi_*.json
Section 8.7 (real fMRI)	lebel_register.py, lebel_sampler.py, lebel_project.py, lebel_features.py, lebel_apply.py, lebel_eval.py	results/ realdata_audio_text.json
Figures	make_figures.py, make_roi_figs.py, make_brain_figs.py, make_topo_grid.py, make_realdata_fig.py	figures/
Demonstration video	demo_run.py, make_demo_video.py	demo/ demo_inversion.mp4

Four Colab notebooks in the notebooks folder run on a free T4 GPU, step by step: Demo_and_Synthetic.ipynb (all synthetic analyses), RealData_TextAudio.ipynb (the real-fMRI analysis), Live_Demo.ipynb (a ten-minute live inversion for demonstrations) and Verification_Tests.ipynb (the corrected barrier and sensitivity tests of Sections 8.3–8.4).